

Introduction

Johanni Brea

Introduction to Machine Learning

Table of Contents

1. What is Machine Learning?

2. Examples

3. Machine Learning: Tasks and Goals

4. Quiz

What is Machine Learning?

- ▶ The term “Machine Learning” appears in the 1960s in the field of Artificial Intelligence (Computer Science), but important concepts existed already before.
Samuel, A. (1959). "Some Studies in Machine Learning Using the Game of Checkers".
<https://doi.org/10.1147/Frd.33.0210>, Schmidhuber, J. (2015). Deep learning in neural networks: An overview <http://dx.doi.org/10.1016/j.neunet.2014.09.003>.
- ▶ Machine Learning shares with Statistics the goal of learning from data.
- ▶ Alternative terms: “statistical learning”, “data science”.
- ▶ Machine Learning started to flourish in the 1990s and gained a lot of popularity since 2010.
- ▶ Some say we are now in the Third AI Summer <https://vimeo.com/389560858>
- ▶ Machine Learning searches for artificially intelligent algorithms that improve with experience: “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P if its performance at tasks in T, as measured by P, improves with experience E.” Mitchell, T. (1997). Machine Learning.

Table of Contents

1. What is Machine Learning?

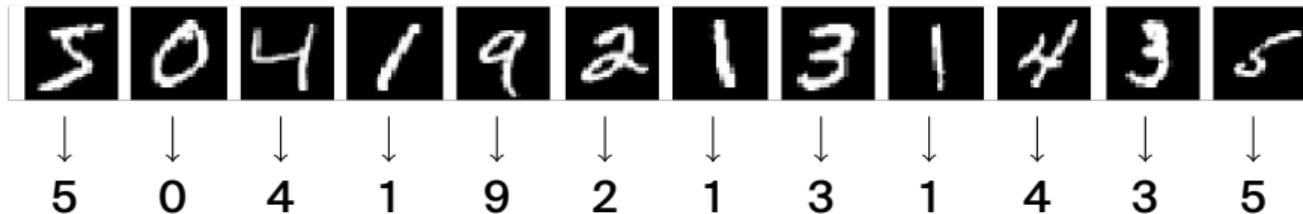
2. Examples

3. Machine Learning: Tasks and Goals

4. Quiz

Examples: Recognition of Handwritten Digits

input: gray scale images (28 x 28 pixels)



desired output: digit class

learning task: learn to imitate humans,
given 60'000 examples of images with labels (i.e. corresponding digit class).

<https://www.kaggle.com/c/digit-recognizer/overview>

Examples: Object Recognition

input: color images



mite container ship motor scooter leopard grille mushroom cherry Madagascar cat
desired output: class label

learning task: learn to imitate humans,
given many examples of images with labels (i.e. corresponding object class).

<https://image-net.org>, <https://devopedia.org/imagenet>

Where Do Labels Come From?

amazonmechanicalturk
Artificial Artificial Intelligence

Your Account HITS Qualifications

Xiaodan Zhou | Account Settings | Sign Out | Help

Introduction | Dashboard | Status | Account Settings

Mechanical Turk is a marketplace for work.
We give businesses and developers access to an on-demand, scalable workforce.
Workers select from thousands of tasks and work whenever it's convenient.

264,053 HITs available. [View them now.](#)

Make Money
by working on HITs

HITs - Human Intelligence Tasks - are individual tasks that you work on. [Find HITs now.](#)

As a Mechanical Turk Worker you:

- Can work from home
- Choose your own work hours
- Get paid for doing good work

Find an interesting task → Work → Earn money

Find HITs Now

or learn more about being a Worker

Get Results
from Mechanical Turk Workers

Ask workers to complete HITs - Human Intelligence Tasks - and get results using Mechanical Turk. [Register Now.](#)

As a Mechanical Turk Requester you:

- Have access to a global, on-demand, 24 x 7 workforce
- Get thousands of HITs completed in minutes
- Pay only when you're satisfied with the results

Find your account → Load your tasks → Get results

Get Started

FAQ | Contact Us | Careers at Amazon | Developers | Press | Policies | Blog
©2005-2012 Amazon.com, Inc. or its Affiliates

An [amazon.com](#) company

<https://ghostwork.info/>

Examples: Cancer Detection

input: images



desired output:

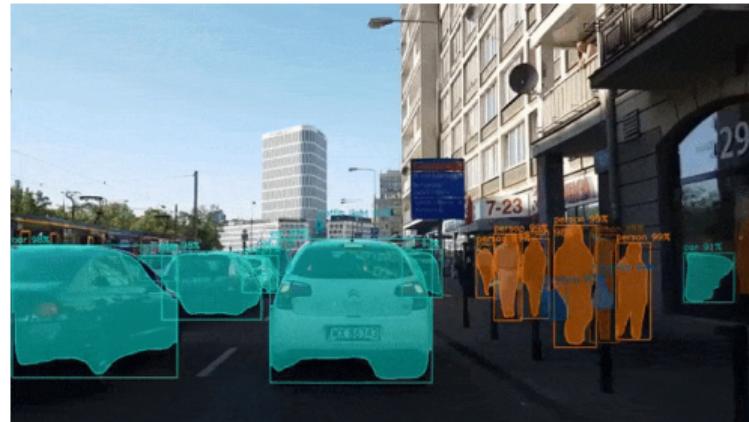
yes, if there is at least one pixel of tumor tissue on the image, otherwise **no**.

learning task: learn to imitate doctors,
given many examples of images with labels (i.e. tumor yes or no).

<https://www.kaggle.com/c/histopathologic-cancer-detection/overview>

Examples: Image Segmentation

input: color images or videos



desired output: assign pixels to objects

learning task: learn to imitate humans,
given many examples of images with labels for every pixels.

https://github.com/matterport/Mask_RCNN, <https://github.com/facebookresearch/detectron2>

Examples: Distance Estimation

input: videos from 8 cameras in a Tesla car (views from 3 cameras shown here)



desired output: positions of and distances to neighboring objects (cars)

learning task: learn to imitate radar data,
given many examples of videos with labels for neighboring objects.

<https://youtu.be/g6b0wQdCJrc>

Examples: Pose Estimation

input: color images or videos



desired output: coordinates of body parts

learning task: learn to imitate humans,
given many examples of images with labels (i.e. corresponding coordinates).

<http://www.mousemotorlab.org/deeplabcut/>

Examples: Image Generation

optional input: high level aspects, e.g. hair color, facial expression, gender, age



desired output: realistically looking images

learning task: learn to generate realistic photos, given many examples of photos.

<https://github.com/NVlabs/stylegan>, <https://thispersondoesnotexist.com/>

Examples: Style Transfer



input: one content image (top left) and one style image (small images)

desired output: an image with desired content and style

learning task: there is no classical learning task here, but rather a creative use of a machine learning model (convolutional neural network).

<https://arxiv.org/abs/1508.06576v2>

Examples: Text Generation

optional input: e.g. start of phrase, context.

Machine learning is not that new. As early as 2006, one team found a way to automatically analyze the social media profiles of millions of people to determine the personalities of them, including the people who liked and commented on the content of a particular post. The technology was so impressive that other companies started experimenting with similar approaches , including a company called DeepMind that won a Google award for its work. But there is a serious problem here , and it is the fundamental lack of interest among academics and the media in solving the problem of social media data in general.

desired output: sensible text

learning task: learn to generate sensible text, given many examples of texts.

<https://github.com/huggingface/transformers>

Examples: Machine Translation

Parteichefin Kipping warnt vor Abschottung - und kritisiert indirekt ihre Genossin Wagenknecht. Die Linken-Bundesvorsitzende Katja Kipping hat ihre Partei aufgefordert, sich in der Frage zwischen Abschottung oder Solidarität eindeutig zu positionieren. Zum Auftakt des Europaparteitags in Bonn sagte Kipping: "Im Sinne dieser Eindeutigkeit sage ich auch: Europa ist längst ein Kontinent der Einwanderung." Das "Recht auf weltweite Bewegungs- und Auswanderungsfreiheit" sei "ein hohes Gut", Menschenrechte unteilbar. Indirekt widersprach Kipping damit auch der Fraktionsvorsitzenden Sahra Wagenknecht, die immer wieder Vorbehalte gegen eine zu flüchtlingsfreundliche Politik geäußert und deshalb auch Kanzlerin Angela Merkel (CDU) mehrfach kritisiert hatte. Wagenknecht ist als bekannteste Politikerin der Linken bereits seit mehreren Wochen erkrankt. Sie nimmt deshalb auch nicht an dem Parteitag in Bonn teil. Der Parteitag spendete ihr Genesungswünsche. Die Bundestagsfraktion hatte einen gegen sie wegen der Kontroversen um die Flüchtlingspolitik geplanten Aufstand Anfang Januar abgeblasen. Kipping sagte vor den Delegierten in Bonn weiter: "Während Trump, Salvini, Orban und Typen wie Seehofer eine Internationale der Mauerbauer schmieden, setzen wir auf internationale Solidarität." Dies sei für die Linkspartei "auch eine Lehre aus der historischen Erfahrung eines Staatssozialismus, der glaubte, mit Mauern und Stacheldraht überleben zu können". Auf eine andere EU hinzuarbeiten, sei die größere Liebeserklärung an Europa als zuzulassen, dass die EU bleibe wie sie ist, erklärte die Parteivorsitzende. Denn der jetzige Zustand spielt den Rechten und Marktradicikalen in die Hände.

La présidente de gauche Katja Kipping a demandé à son parti de se positionner clairement sur la question du cloisonnement ou de la solidarité. En prélude à la journée de l'appartement européen à Bonn, Kipping a déclaré : "Dans le sens de cette clarté, j'ai également dit que l'Europe était depuis longtemps un continent d'immigration". Le "droit à la liberté de circulation et d'émigration dans le monde" est "un bien précieux", les droits de l'homme indivisibles. Kipping contredit ainsi indirectement le président du groupe Sahra Wagenknecht, qui avait toujours exprimé des réserves sur une politique trop favorable aux réfugiés et avait donc critiqué à plusieurs reprises la chancelière Angela Merkel (CDU). Wagenknecht, la politicienne la plus connue de la gauche, est malade depuis plusieurs semaines. Elle ne participe donc pas non plus au congrès du parti à Bonn. Le congrès du parti lui a adressé ses voeux de rétablissement. Le groupe du Bundestag avait fait sauter un soulèvement prévu contre lui au début du mois de janvier en raison des controverses sur la politique des réfugiés. Kipping a ajouté devant les délégués à Bonn : "Alors que Trump, Salvini, Orban et des types comme Seehofer forgent une Internationale des constructeurs de murs, nous missons sur la solidarité internationale". Pour le parti de gauche, "c'est aussi une leçon de l'expérience historique d'un socialisme d'État qui croyait pouvoir survivre avec des murs et des fils barbelés". Travailler vers une autre UE est une plus grande déclaration d'amour à l'Europe que de permettre à l'UE de rester telle qu'elle est, a déclaré la présidente du parti. Car la situation actuelle fait le jeu de la droite et des radicaux du marché.

<http://www.statmt.org/wmt19/index.html> (Microsoft Research Asia: Multi-Agent Dual Learning)

Examples: Automatic Captions

input: image



desired output: caption

“A small train on a city street with people near by.”

learning task: learn to imitate humans, given many examples of captioned images.

<https://github.com/microsoft/Oscar>

Examples: Speech Synthesis and Speech Recognition

input: written sentence

“The essential point to be remembered is that the ornament, whatever it is, whether picture or pattern-work, should form part of the page,“

desired output: spoken sentence



learning task: learn to imitate humans,
given many examples of text with recordings of the corresponding speech.

<https://nv-adlr.github.io/WaveGlow>

For speech recognition, recordings are given as input and written text is the output
(see e.g. <https://arxiv.org/abs/2105.00982v1>).

Examples: Theorem Proving

input: mathematical axioms and theorems (goals)

cnf(sos01,axiom, (product(A,A) = A)).

cnf(sos02,axiom, (product(A,product(B,C)) = product(product(A,B),product(A,C)))).

...

cnf(goals,negated_conjecture, (product(product(product(x0,x1),x1),product(x0,x2)) != product(product(x0,x1),product(product(x1,x0),x2)))).

desired output: steps to prove the theorem from the axioms

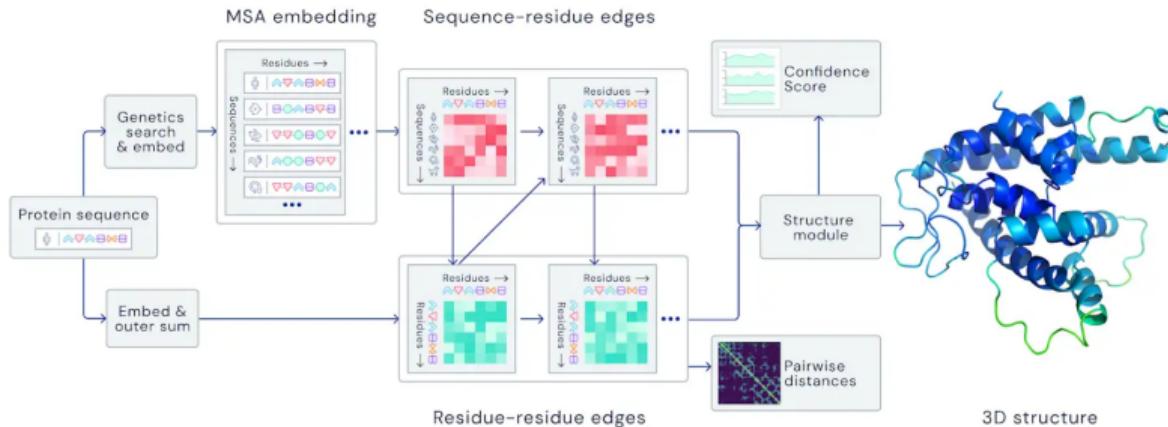
learning task: learn solving strategies by trial-and-error on many examples of axioms and theorems.

<http://www.tptp.org/>, [https:](https://)

//deepmind.com/research/publications/Training-a-First-Order-Theorem-Prover-from-Synthetic-Data

Examples: Protein Folding

input: protein amino acid sequence



desired output: 3D structure of the protein

learning task: learn to imitate the laborious task of 3D reconstruction (X-ray crystallography) from many examples of protein sequences with known 3D structure.

<https://deepmind.com/research/case-studies/alphafold>

Examples: Machine Learning as Models of Brain Function

Questions & Answers | Open Access | Published: 15 August 2011

Machine learning for neuroscience

Geoffrey E Hinton 

[Neural Systems & Circuits](#) 1, Article number: 12 (2011) | [Cite this article](#)

A Neural Substrate of Prediction and Reward

Wolfram Schultz, Peter Dayan, P. Read Montague*

* See all authors and affiliations

Science 14 Mar 1997;

Prefrontal cortex as a meta-reinforcement learning system

Jane X. Wang, Zeb Kurth-Nelson, Dharshan Kumaran, Dhruva Tirumala, Hubert Soyer, Joel Z. Leibo, Demis Hassabis & Matthew Botvinick 

[Nature Neuroscience](#) 21, 860–868 (2018) | [Cite this article](#)

Using goal-driven deep learning models to understand sensory cortex

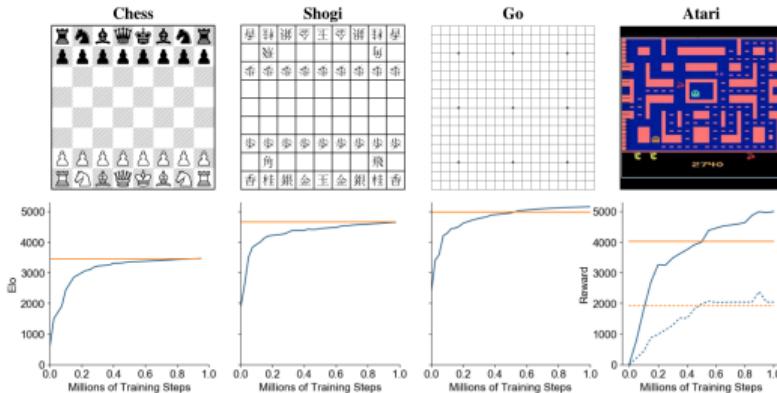
Daniel L K Yamins  & James J DiCarlo

[Nature Neuroscience](#) 19, 356–365 (2016) | [Cite this article](#)

- ▶ <https://neuralsystemsandcircuits.biomedcentral.com/articles/10.1186/2042-1001-1-12>
- ▶ <https://www.pdn.cam.ac.uk/system/files/documents/1997-science.pdf>
- ▶ <https://www.biorxiv.org/content/biorxiv/early/2018/04/13/295964.full.pdf>
- ▶ <http://brainmind.umin.jp/PDF/wt17/Yamins3.pdf>

Examples: Games

input: rules of a game



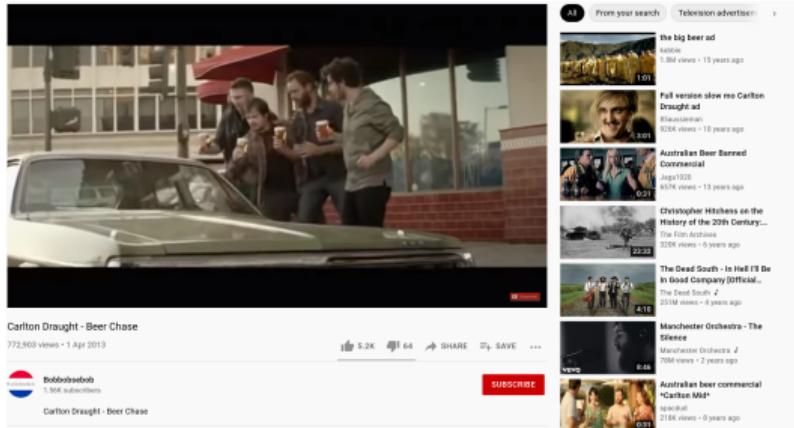
desired output: winning policy

learning task: learn winning strategies by trial-and-error from many games of the computer playing against itself.

<https://arxiv.org/abs/1911.08265>

Examples: Advertisement

input: user profile (currently viewed page & history)



desired output: attractive suggestions

learning task: learn by trial-and-error to display the suggestions that users will select with high probability.

Other Links

Competitions

<https://www.kaggle.com/>

<https://dreamchallenges.org/>

Notable Companies

[https://www.amazon.science/](https://www.amazon.science)

<https://ai.facebook.com/>

<https://ai.google/>

<https://deepmind.com/>

<https://machinelearning.apple.com/>

<https://openai.com/>

Datasets

<https://www.datasetlist.com/>

<https://registry.opendata.aws/>

<https://www.kaggle.com/>

<https://www.openml.org/>

State-of-the-Art Research

<https://paperswithcode.com/sota>

Books

<https://people.eecs.berkeley.edu/~russell/hc.html>

https://en.wikipedia.org/wiki/The_Age_of_Surveillance_Capitalism

<https://ghostwork.info/>

Table of Contents

1. What is Machine Learning?

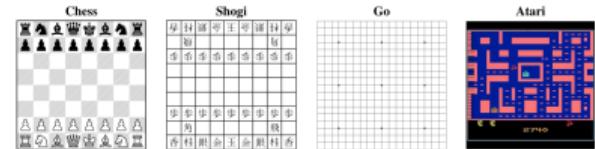
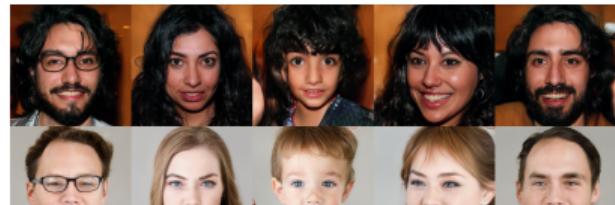
2. Examples

3. Machine Learning: Tasks and Goals

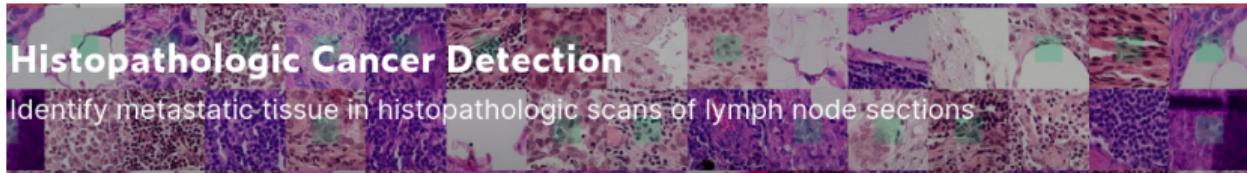
4. Quiz

Tasks

- ▶ **supervised learning:** predict a class label (**classification**) or real values (**regression**) from an input. Training data: input X , output y .
- ▶ **unsupervised learning:** discover good features for representing or visualizing the input data or for generating examples similar to the input data. Training data: input X .
- ▶ **reinforcement learning:** discover what action should be performed next (a policy) in order to maximize the eventual payoff. Training data is obtained by interacting with an environment and observing reward.



Goals



- ▶ **Prediction** with highest possible accuracy.
e.g. "Do we see cancer on this image or not?"
- ▶ **Interpretation**.
e.g. "Which visual features are most indicative of cancer?"
- ▶ **Visualization**.
e.g. "Are some images more similar to each other than to others, i.e. are there some clusters in the data?"
- ▶ **Data Generation**.
e.g. "Can we generate realistically looking artificial histology images?"

Table of Contents

1. What is Machine Learning?

2. Examples

3. Machine Learning: Tasks and Goals

4. Quiz

Quiz

1. What kind of machine learning task is the recognition of handwritten digits?
A supervised **B** unsupervised **C** reinforcement
2. What kind of machine learning task is machine translation?
A supervised **B** unsupervised **C** reinforcement
3. What kind of machine learning task is speech synthesis?
A regression **B** classification **C** neither
4. What is the primary goal in the above three tasks?
A prediction **B** interpretation **C** visualization
5. What kind of machine learning task is theorem proving?
A supervised **B** unsupervised **C** reinforcement
6. What kind of machine learning task is cancer detection?
A regression **B** classification **C** neither

Programming in Julia

- ▶ Code examples in Pluto notebooks and on <https://bio322.epfl.ch>.
- ▶ Installation instructions on <https://github.com/jbrea/MLCourse>.
- ▶ Recommended Workflows
 1. Write your solutions directly into the notebooks.
 2. Create your own Pluto notebooks for the solutions.
 3. Use VS Code or some other editor
<https://github.com/julia-vscode/julia-vscode>.
 4. If you choose option 2 or 3 you should first activate the MLCourse environment (in option 1 this is done already in a hidden cell):
`using Pkg; Pkg.activate(joinpath(Pkg.devdir(), "MLCourse"))`
- ▶ Get help from the Julia community <https://julialang.org/community/>.

Recommended Textbooks

This course is strongly inspired by the first two books.
The other books are more advanced.

- ▶ **An Introduction to Statistical Learning**

www.statlearning.com

- ▶ **Reinforcement Learning an Introduction**

<http://incompleteideas.net/book/the-book-2nd.html>

- ▶ **The Elements of Statistical Learning**

<https://web.stanford.edu/~hastie/ElemStatLearn>

- ▶ **Bayesian Reasoning and Machine Learning**

<http://web4.cs.ucl.ac.uk/staff/D.Barber/pmwiki/pmwiki.php?n=Brml.HomePage>

- ▶ **Pattern Recognition and Machine Learning**

www.microsoft.com/en-us/research/publication/pattern-recognition-machine-learning